

Amazon SageMaker Components (Console Order)

Applications & IDEs

Includes Studio (IDE), Notebooks (Jupyter), Canvas (no-code ML), and RStudio. Used to write code, explore data, or build ML models without coding.

Admin Configurations

Settings for managing user roles, permissions, and policies. Ensures correct access control to SageMaker resources.

JumpStart

Collection of prebuilt ML solutions and pre-trained models. Helps to quickly start ML projects without coding everything from scratch.

Governance

Provides compliance, lineage tracking, and auditing tools. Ensures responsible AI usage and enterprise-level security requirements.

HyperPod Clusters

Specialized compute clusters for large-scale distributed training. Useful for training huge models like LLMs with optimized GPU clusters.

Ground Truth

Managed data labeling service for creating high-quality training datasets with human and automated labeling options.

Processing

Run preprocessing and postprocessing jobs at scale. Prepares data before training and evaluates results after training.

Training

Service to train ML models using built-in or custom algorithms. Supports distributed training for scalability.

Training Plans

Helps organize and track multiple training jobs. Useful for managing experiments and comparing model runs.

Inference

Deploy trained models for predictions. Supports real-time endpoints, batch inference, and serverless inference.

Augmented AI (A2I)

Combines AI model predictions with human review. Useful in compliance-heavy or high-risk use cases.

AWS Marketplace

Marketplace for third-party ML models, datasets, and algorithms. Enables 'buy instead of build' for faster adoption.